

Inference Latency: FP32 vs INT8

